

Why I'm leaving OpenAI to build telepathy

August 4, 2026

On Thursday, July 23rd, I resigned from OpenAI. On the 24th, I started as a Founding Researcher at [Conduit](#). We're building telepathy: thought-to-text models, trained on non-invasive neural data.

I'll talk about:

1. **Thought-to-text**: what using telepathic tools will be like in 2027, 2030, and 2035
2. **How do we build thought-to-text**: why Conduit is collecting immense quantities of non-invasive neural data
3. **Why I joined**: I joined because of the audacious vision, the brilliant people, and the fun problems
4. **Hi**: about me

1. Thought-to-text, aka telepathy

My prediction: In a couple years, the main way we'll talk with our AIs is with our thoughts. They will not be [a super-smart automated intern or coworker](#) that you desperately try to keep up with. They will be a natural, joyful extension of you.

I wrote down vignettes of this future below. I view them as optimistic but highly plausible.

2027 — A morning at Conduit

I put a band around my head, and open my laptop at Conduit. It's 9 am. My device pairs with my laptop over Bluetooth. I open, without loss of generality, Codex.

I think about the day's work as I look at my code, review yesterday's notes, and prepare for our 9:30 am standup. My GPT-7 agent has been chugging away at the new encoder I've been exploring for over a day now. I look at its work. I become confused by the plot labels, annoyed at the AI-speak in its first paragraph, and curious about the symbols in the equations. My vague thoughts get sent to Conduit's model, which uses its priors over language and the kinds of things I might say to output:

```
conduit:///thought-stream → codex
```

```
$ decode --source=neural-latent --autosend --to=codex
```

```
thought_01 > I'm looking at your plots. Please relabel the legend, and add a text cell  
              below with a simple description of each line.
```

```
thought_02 > I understand your 2nd paragraph of discussion, but please rewrite the first  
              paragraph using the humanizer skill.
```

```
thought_03 > Make a textbox explaining the equations with cleaner notation.
```

```
streaming >
```

The plot, paragraphs, and notation each took me 10 seconds to glance over. I have auto-send on, so after each 10-second chunk, the Conduit model deciphers my thoughts and sends it to GPT-7, which assigns each incoming task to a subagent. I look at the new plot, paragraph, and equations. I wonder whether spherical harmonics might be useful after all; Codex spins off a subagent. Convinced by the plot, I decide to make a quick slideshow with a prettified version showing only the baseline and top two lines, to show my teammates; Codex spins off a subagent.

I make myself coffee. My thoughts are mostly empty, but I briefly recall what I want to say to the candidate I'm getting lunch with tomorrow. In the background, Codex is prompted to think about info I might want for the lunch chat, and decides that updating the synthetic data scaling plot before showing it off would be prudent; Codex spins off a subagent.



I'm not saying words **really loudly** in my head while getting coffee. I just read the plots as I normally do, and make coffee as I normally do. It feels like magic.

2030 — Industry adapts and Conduit expands

The AI companies now train their models to directly interface with Conduit's latent representations. My encoded thoughts get sent directly to Codex, rather than having to pass through Conduit's decoder model first. That means I easily communicate thoughts that are hard to describe in text, like mental images.

I rarely go band-less when chatting with AIs these days. It's annoying, honestly, to go without. With the neural headband, I feel like I have superhuman powers over my laptop! Without it, I feel like I'm talking with a superpowered alien who's trying its best to be helpful but isn't sure what I want and is scared I'll get mad if it does something I don't want. Ugh.

Conduit continues to iterate on non-invasive read, but we've spun up two new efforts: invasive general read, and general write.

+-----+ ASIDE: THREE AXES TO NEUROTECH +-----+			
NON-INVASIVE	<-----	ACCESS	-----> INVASIVE
no surgery			requires surgery
+-----+			
READ	<-----	DIRECTION	-----> WRITE
extract information from brain:			make changes to the brain:
intentions, actions, emotions,			hear audio, or tell the brain
internal monologue, memories			to use more calories
+-----+			
NARROW	<-----	SCOPE	-----> GENERAL
one predefined capability			many capabilities from one interface
+-----+			

Most people are happy to stick with their neural bands, but a good number are excited to get higher fidelity reads via invasive tech.

More interesting is the recent excitement in writes. I feel superhuman in my control over my laptop. But my senses are still merely human. It's like if I could control my arms, but I'd lost all feeling in them. Yeah, I can still *see* my arms, and it's way better than not having arms, but it's still really odd. Like [Ian Waterman](#). I want to *feel* what my Codex feels. Now that Conduit does general read, learning to do general write is many OOMs more data efficient.

More importantly, I want to unlock the other applications of write technology. I want to make my brain as efficient and neuroplastic as when I was 18.

2035 — The AI is no longer “[other](#)”

My AI is a natural extension of me. It feels like a sixth sense and another limb. I wonder about a problem, feel as though I'm literally surfing the web, see glimpses of the websites, get flashes of intuition about the problem, and ultimately derive the answer. It feels fun! My brain is like the Flash. Because of write tech, my thinking is the fastest it's ever been even when I turn my AI off.

How superhuman 2035 looks depends directly on how superhuman we, as society, decide to make our AIs. Perhaps we choose to [pace ourselves](#). But I like that this is literally a human-in-the-loop vision of the future, where AI directly empowers humans rather than replacing us.

► A few quotes about the future of thought-to-AI

2. How do we build thought-to-text?

In theory, it's simple. Our input is brain activity, and our target output is what the person was doing at the time – for example, what text the person wrote. Given the brain activity,

we want to predict output that is semantically similar to what the person wrote.

To train models that can predict text given brain signals, we must apply the same lesson learned by those predicting text given speech audio, or text given preceding text: the [bitter lesson](#). The lesson roughly states that you should throw more useful compute at your model, and your model will become better than any ingenious algorithm you could've hand-crafted. That means we must scale up our data collection by orders of magnitude beyond what has ever been done in academia.

It's hard to collect enough data using invasive methods. Few people want a chip in their brain! But non-invasive methods are getting much better. The hardware is improving and getting cheaper, though I apologize for being vague about the particularities of our hardware.

As we're training on more data, the model is predicting text that is more semantically similar to the subject-written text. Yes, there's some irreducible error due to noise, but for most modalities we're not yet in a regime where we're pushing against that. Concretely, the scaling laws are looking good: the cosine similarity of our latent space predictions with the target latent spaces goes up as a straight line with respect to the logarithm of the number of hours of data. We're in the GPT-2 era.

We don't need perfect decoding to be useful. Your thoughts will be like GPS in a city: a noisy GPS signal isn't enough to determine your exact location. But combined with a map and a navigation route — equivalently, the LLM and context — it becomes remarkably accurate.

To read up on how we do data collection, check out Conduit's [blog post](#). If you're more of the active learning type, come be a [research participant](#)!

+-----+ IMPERFECT DECODING CAN STILL BE INCREDIBLY USEFUL +-----+				
YOUR THOUGHTS	+	LLM + CONTEXT	=	PRECISE INFERENCE
= NOISY GPS SIGNAL	+	= MAP + ROUTE	=	
? *	+	[---o-----]	=	[---o--*--] -> answer
*	+	[~~~]	=	[~~~]
+-----+				

3. Why I joined

3 reasons:

1. **Audacious yet practical vision:** I believe in this vision; to quantify, if Conduit becomes the general read and write company, then I expect it to be worth over \$1 trillion. I feel good about the vision, because it builds towards a more human future. And I trust the scaling laws, which have held for many doublings.
2. **Brilliant people:** The co-founders, Rio Popper and Clem von Stengel, are brilliant and relentless and sweet. I love working with them, and I love working with the several others who've joined to work on this same mission, on the same intellectually interesting problem.
3. **Fun problems:** The research is SO interesting. I get to think about encoders, data efficiency, synthetic data, data attribution, multimodality, weird statistics... it's a blast. At OpenAI, you can only work on a narrow slice of The Problem, and even on that narrow slice you're constrained by the existing architecture. At Conduit, it's all greenfield.

In short: I'm happy! I'm working on a problem I'm obsessed with, that is frighteningly ambitious, with a small group of people I like.

4. By the way, hi!

For those who don't know me: hello! I left OpenAI two weeks ago, after spending 1.5 years there as a researcher.

I grew up in Washington State, where I played competitive chess from ages 5–15 and stopped after becoming a [WIM](#). My first time hearing about the potential for smarter-than-human AI was when I was 12, but for years I figured it was just a weird but interesting idea that people on the internet liked to write about.

At 18, I enrolled in Harvard to study computer science. In a class, I learned about GPT-3 and finally read Bostrom's *Superintelligence*. Wow, what a wakeup moment. I got invested in AI safety research, had a brief stint in AI policy, and ultimately joined the alignment team at OpenAI. At OpenAI, other than my research, I spent some time on various side projects, including making OpenAI's AGI onboarding presentation and our [alignment blog](#), and helping advise the AI Resilience division of the OpenAI Foundation.

If you want to chat about Conduit, reach out to me at naomi@condu.it. Let's grab coffee, or I'll give you a tour of our unusual and beautiful office in San Francisco. We're always hiring researchers, infra folks, and operators.

Acknowledgements: Thank you to Aidan Smith, Devansh Pandey, Julia Shephard, Lev Chizhov, Ryan Kaufman, and the Conduit team for feedback. I'm grateful, too, for the several other companies building ambitious BCI tech, including but not limited to those building better sleep, better ultrasound, a solution for blindness, and uploads - I'm excited to use your tech! All mistakes are my own.

[← Back to Blog](#)

[LinkedIn](#) [X \(Twitter\)](#) [Email](#) [RSS](#)